

From Erasure to Obfuscation

A Paradigm Shift in Privacy in the Age of Generative AI

Raymond

First draft: May 2026

© Raymond 2026

This work is licensed under Creative Commons Attribution-NonCommercial 4.0 International (CC BY-NC 4.0).

You may share and adapt this work for non-commercial purposes with attribution.

Full license: <https://creativecommons.org/licenses/by-nc/4.0/>

AI Collaboration Disclosure

This essay was developed in collaboration with Claude (Anthropic). The conceptual direction, judgments, structural decisions, and final voice belong to the author; Claude assisted with literature integration, drafting, and language polishing. All factual claims have been independently verified.

Suggested citation:

Raymond (2026). From Erasure to Obfuscation: A Paradigm Shift in Privacy in the Age of Generative AI. Working paper.

Abstract

The "erasure paradigm" dominating privacy governance since the rise of the internet—exemplified by Article 17 of the GDPR and Article 47 of China's Personal Information Protection Law—faces theoretical limitations as generative AI reshapes the infrastructure through which individuals are publicly known. As large language models increasingly mediate access to information about persons, the central privacy threat is shifting from "data being recorded" to "data being automatically integrated into stable profiles." This essay argues that traditional erasure pathways are neither sufficient nor uniquely necessary; privacy protection may need a complementary strategy that weakens profile stability through the deliberate production of internal inconsistency in self-representation. The essay names this strategy *narrative autonomy*, situates it within and distinguishes it from five adjacent traditions—presentation of self, context collapse, obscurity, contextual integrity, and obfuscation—and identifies its distinctive contribution: resistance directed at automated integration rather than at data collection or audience management. The essay further offers a speculative observation that model collapse may carry an under-recognised privacy function. The contribution is conceptual.

Keywords: generative AI; personal information protection; right to be forgotten; model collapse; narrative autonomy; obfuscation

1. Introduction: The Limits of Erasure

For roughly two decades, mainstream privacy governance has rested on a deceptively simple assumption: that the threat to privacy comes from the existence of data, and that the path to protection is therefore to make data cease to exist. This assumption is embedded deep in contemporary law. Article 17 of the GDPR establishes a "right to erasure (right to be forgotten)" under which data subjects can require controllers to delete their personal data in specified conditions (European Parliament and Council, 2016). Article 47 of China's PIPL, enacted in 2021, similarly stipulates that personal information processors shall delete personal information when the processing purpose has been fulfilled, when the retention period has expired, when consent is withdrawn, or under several other conditions; if processors do not delete it, individuals have the right to request deletion (Standing Committee of the National People's Congress, 2021).

Yet the erasure paradigm has always faced a structural problem. The internet's replication mechanisms ensure that any once-public information may persist in mirrors, backups, archives such as the Wayback Machine, and screenshots. Legal "deletion" rarely corresponds to physical disappearance. Brunton and Nissenbaum (2015), in their book-length treatment of *obfuscation*, observe that in distributed data ecosystems, deletion alone is structurally insufficient, and they propose that obfuscation—the deliberate production of misleading, ambiguous, or confusing information to interfere with data collection—should be added to the privacy toolkit. The literature on the right to be

forgotten has, in parallel, traced a long line of execution failures, jurisdictional conflicts, and free-expression frictions (Jones, 2016; Mantelero, 2013).

This essay extends those concerns, but identifies a development that the existing literature has not yet fully theorised. Both the erasure paradigm and the obfuscation tradition share an unstated premise: that the primary readers of personal information are humans, or human-operated systems for targeted advertising and surveillance. Under that premise, deleting a social-media post largely removes it from human attention, and obfuscating one's search history largely thwarts behavioural advertising. But this premise is becoming inaccurate. Increasingly, the most consequential reader of distributed personal information online is a generative AI system that integrates fragments into a single, fluent, authoritative-seeming representation of the person.

This shift is more than a technical curiosity. It restructures what privacy threats look like, what governance instruments can address them, and what individual strategies of resistance become available. The remainder of this essay develops three claims. **First**, the infrastructural shift from human to machine reading constitutes a paradigm change for privacy governance, not merely an extension of existing conditions. **Second**, this shift recasts the privacy threat: the central concern is no longer the existence of data but the *statistical coherence* with which data is integrated into automated profiles. **Third**, this re-framing creates space for a complementary privacy strategy I call *narrative autonomy*: a privacy interest in publishing multiple, internally inconsistent self-narratives in order to weaken algorithmic profile stability, without necessarily deleting any data.

The contribution of this essay is conceptual rather than empirical. The argument is offered as a re-framing that opens new questions for privacy theory and policy. The claim is not that erasure governance should be abandoned, but that a category of privacy interests has emerged which erasure alone cannot serve.

2. Infrastructural Shift: From Human to Machine Reading

Datafication, as theorised by van Dijck (2014), refers to the rendering of social life into machine-readable data, treated as a quasi-natural resource. Critical data studies has charted, in detail, how this transformation reshapes power relations (Couldry & Mejias, 2019; Zuboff, 2019), epistemic authority (Pasquale, 2015), and the conditions under which selfhood is constituted in public (boyd & Crawford, 2012). Yet much of this literature was written before generative AI matured into a mass-consumer interface. The infrastructural change of the past three years—roughly from the public release of ChatGPT in late 2022 onwards—warrants its own analysis, because it changes both *who* reads personal information online and *how* that reading is done.

Empirically, the scale of adoption is substantial. Pew Research Center's nationally representative survey conducted in early 2025 found that 34% of U.S. adults reported having used ChatGPT, roughly

double the 2023 share, with adoption rising to 58% among adults under 30 (McClain & Sidoti, 2025). The Pew survey did not isolate "looking up information about other people" as a discrete use category. What is documented is rapid use of generative systems for tasks that previously fell to search engines, and the integration of generative interfaces into mainstream search products (Google AI Overviews, Microsoft Copilot in Bing) and productivity software. The argument here does not require generative AI to be the dominant person-search interface today; it requires only that it become a sufficiently consequential one that its mediating properties begin to shape privacy outcomes.

Three properties of this new mediating layer matter for privacy.

First, model "reading" is statistical. A large language model does not "look up" facts about an individual the way a human reading a profile page would. It generates outputs by probabilistic inference over patterns abstracted from training data and, where retrieval-augmented generation is enabled, from documents pulled at inference time. A person's representation in a model is, in effect, the statistical compression of every fragment of text the model has encountered about them. This is a different epistemic object from a search result list. Where a search result preserves the heterogeneity and provenance of sources, a generative response collapses heterogeneity into apparent coherence.

Second, model output carries an authority effect. Where a search result list visibly aggregates many sources, a model's prose response presents itself as a single, coherent answer. Studies of trust calibration suggest users systematically over-trust fluent generative output (Bender et al., 2021; Weidinger et al., 2022). This creates an asymmetry: the model's representation of an individual carries more apparent authority than the underlying data warrants. Pasquale's (2015) "black box society" applies here with new force, because the generative interface adds *fluency* to opacity.

Third, the model becomes a new bottleneck of public knowledge. Search engines exposed users to plurality—competing pages, conflicting accounts, unresolved tensions. Generative interfaces compress these into a single response. The narrative about a person becomes less plural, more singular, and more dependent on whatever the model has converged on. This compression matters because it removes the friction by which contradictory accounts of a person could resist consolidation.

Taken together, these properties shift the locus of privacy concern. In the search-engine era, the worry was that data, once published, was permanently retrievable. In the LLM era, the worry has a different shape: that data, once published, will be automatically *integrated* into a stable representation that has the authority of a single answer. This is not a continuation of the old privacy threat with new technology; it is a different threat, calling for different governance.

3. From Memory to Statistical Coherence

In the older framing, the threat chain ran: information exists → a human encounters it → the individual is affected. Severing this chain required making information not exist—hence erasure.

In the new framing, the threat chain runs: information exists → a model integrates it into a profile → that profile drives decisions in hiring, credit, insurance, social judgment, and increasingly automated processes (Eubanks, 2018; O’Neil, 2016) → the individual is affected. The critical link is not the existence of data but the success of integration. If the integration step fails—if the model cannot reliably converge on a stable representation of the person—the threat does not materialise even if every original fragment of data still exists.

This is the analytic move worth foregrounding: **the central privacy threat is no longer memory but statistical coherence**. What harms a person is not that data about them exists somewhere, but that an automated system can stitch that data into a confident profile.

It is worth being precise about what "statistical coherence" means here: the property by which a generative system, given a query about a person, returns outputs that are internally consistent across runs, semantically dense, and presented with high apparent confidence. This is distinct from *accuracy*. A profile may be statistically coherent without being accurate. A coherent but inaccurate profile harms the subject precisely because its consistency is read as reliability. Conversely, a profile that fragments into contradictory outputs across runs—even if individual outputs would have been accurate—fails to function as a basis for downstream decision-making, because the system itself signals it cannot be relied upon.

This re-framing has three consequences. It loosens the strict requirement for erasure, identifies "profile stability" as a measurable governance object, and opens space for individual resistance—which the next section develops.

4. Narrative Autonomy: A Theoretical Sketch

I propose the term *narrative autonomy* to describe a privacy strategy fitted to this new threat structure. The core claim: in environments where automated profiling depends on the statistical coherence of distributed personal data, individuals have a privacy interest in producing internal inconsistency in that data—and a corresponding latitude to do so within limits set by other legal duties (against fraud, defamation of others, and so on).

Concretely, narrative autonomy describes the practice of publishing, across multiple platforms, several plausible but mutually incompatible self-accounts. Different occupations on different bios. Different home cities in different posts. A range of past affiliations that do not all cohere. None of these statements need be flatly false in a way that would constitute fraud against a counterparty in a transaction; they may simply represent different facets, different time-slices, different self-presentations, or genuinely fictionalised accounts of one’s life. The aggregate effect on an LLM

trained or retrieval-augmented over this material is a low-confidence representation, often expressed as hedging, refusal, or overt inconsistency in model output.

4.1 Continuity with Goffman

Goffman's (1959) classic treatment of the *presentation of self in everyday life* established that human social existence routinely involves the strategic enactment of different selves before different audiences. Multiplicity of self-presentation is not pathological; it is constitutive of social life. Narrative autonomy draws on this insight but extends it. Goffman's analysis assumes audience-specific contexts in which each presentation is locally appropriate and sincere. The challenge that arises in the digital age is that the structures of context that made Goffman's framework work are increasingly absent. There is no longer a kitchen and a dining room; there is one platform that addresses everyone, and one model that reads everything.

4.2 Continuity with Context Collapse

Marwick and boyd (2011) gave this problem its now-canonical name. *Context collapse* refers to the flattening of multiple distinct audiences into a single addressable public. Narrative autonomy can be read as a response to a *new layer* of context collapse: an algorithmic layer in which the collapsing agent is not the audience but the model. The earlier context collapse left individuals with the choice of self-presentation but flattened their audiences; the new context collapse takes self-presentations that may have been locally appropriate across different platforms and times, and integrates them into a single algorithmic representation that flattens them as well. Narrative autonomy is the strategy of producing self-presentations that resist this second-order flattening.

4.3 Continuity with Obscurity, Contextual Integrity, and Obfuscation

Hartzog and Stutzman (2013) proposed *obscurity* as a privacy concept oriented to the costs of finding and aggregating personal information. Narrative autonomy is, in one sense, an obscurity strategy fitted to a new architecture: where Hartzog and Stutzman analysed how technical and social friction protected aggregation costs, narrative autonomy targets the integration *quality* once aggregation has occurred. The mechanism is different: obscurity protects through cost; narrative autonomy protects through inconsistency.

Nissenbaum's (2010) framework of *contextual integrity* anchors privacy claims in the appropriate norms of information flow within specific social contexts. Narrative autonomy extends this to an environment in which contexts themselves are dissolved by automated integration. The model that aggregates a person's medical-forum posts, their employment-platform bio, and their fiction-writing accounts violates contextual integrity not at the collection step but at the integration step.

Brunton and Nissenbaum (2015) developed *obfuscation* as a tactical response to surveillance and tracking, with examples ranging from radar chaff to TrackMeNot to credit-card pooling. Their cases

target *data collection*. Narrative autonomy targets a different moment: the *integration* of already-collected data into stable profiles. It is, in this sense, an obfuscation strategy adapted to the LLM era.

4.4 What Narrative Autonomy Adds

Set against these adjacent concepts, narrative autonomy adds two things. **First**, it identifies the *integration step* in automated profiling as a distinct site of privacy concern, not assimilable to existing collection-focused or audience-focused frameworks. **Second**, it articulates a corresponding strategy of resistance—the deliberate production of internally inconsistent self-narratives—and argues that this strategy is consistent with existing legal latitude for self-presentation, so long as the boundary against fraud against specific counterparties is preserved.

Importantly, narrative autonomy is not coextensive with anonymity, pseudonymity, or false identity. It operates *under one's own name*. The individual is not hiding; they are publishing multiple genuine self-presentations that an automated system cannot stably reconcile. This distinguishes the strategy from frameworks already addressed by anonymity doctrine and existing prohibitions on identity fraud.

4.5 Worked Examples

A freelance writer publishes four professional bios across four platforms. On her personal blog she emphasises her literary fiction. On a professional networking site she emphasises her technical writing. On a creative-writing community she discusses an unfinished novel that draws on autobiographical material with deliberate liberties. On a music-review platform she writes under her own name about cultural criticism. Each presentation is sincere within its context; none is intended to deceive any specific counterparty; the aggregate is internally inconsistent. A model attempting to produce a unified profile of her returns hedged or contradictory output. **This is narrative autonomy.**

A graduate student preparing for the academic job market keeps his academic output consistent on his university page and publication record. But he does not curate his older social-media presence to match: a teenage blog with experimental fiction remains under his name, as does an early creative project about an imagined alternate career as a chef. **He does not delete; he leaves the multiplicity in place.** Some hiring committees' AI tools return inconsistent profiles; the human committee, reading carefully, can distinguish his academic work from his older creative work. The model is the only reader for whom the multiplicity is a problem.

A public figure—say, a CEO—publishes inconsistent biographical claims across her speeches, interviews, and personal website to make AI-generated profiles of her unreliable. **This is a misuse of narrative autonomy.** As a public figure with public-interest accountability obligations, her case falls within the public-figure exception developed below.

5. Model Collapse as Accidental Privacy Infrastructure

Shumailov et al. (2024), publishing in *Nature*, demonstrated that when generative models are repeatedly trained on data produced by earlier generative models, they undergo what the authors call *model collapse*: the tails of the original data distribution—rare events, low-frequency entities, distinctive details—are progressively lost, while the bulk of the distribution narrows toward the mean. Subsequent work has both replicated and qualified this finding. Notably, Gerstgrasser et al. (2024) showed that the strong collapse effect arises primarily when synthetic data *replaces* real data at each iteration; when synthetic data is *accumulated alongside* original real data—a setting closer to actual web-scale training conditions—test error remains bounded and collapse is largely avoided. The empirical picture is contested rather than settled.

I propose, as a *speculative observation*, that model collapse may have a second face: an unintended privacy function for low-profile individuals. For those whose representation in training corpora is sparse to begin with, model collapse selectively erodes precisely the kind of long-tail signal that could support a confident profile. Where high-profile individuals would presumably continue to be well-represented—because their information is reinforced through repeated mentions in human-generated, high-quality sources—low-profile individuals may experience an accidental privacy windfall: not because their data has been deleted, but because the integration mechanism is degrading.

This observation is offered as a hypothesis warranting empirical study, not a settled finding. Three conditions would need to hold: (a) generative AI must become a sufficiently dominant mode of information retrieval about persons; (b) recursive training must continue under data-replacement rather than data-accumulation conditions, which Gerstgrasser et al. (2024) showed can avoid collapse; and (c) the long-tail loss observed by Shumailov et al. in benchmark settings must extend to the person-information domain. A further qualification: retrieval-augmented generation may partially compensate for training-data degradation.

If the hypothesis holds, model collapse is not only a technical problem; it is a candidate piece of accidental privacy infrastructure. Proposals to "fix" model collapse by mandating provenance tracking should consider whether such fixes inadvertently restore profiling capacity in ways that ordinary individuals have a privacy interest in resisting. Crucially, the conceptual contribution of the essay does not depend on this hypothesis being correct. Narrative autonomy is a coherent privacy strategy whether or not model collapse provides accidental support for it.

6. Risks and Boundaries

Fraud and reliance interests. A central worry is that narrative autonomy slides into deception in transactional contexts. The proposal is not that all forms of self-narration are protected. Existing legal duties—against fraud in contractual relationships, against material misrepresentation in regulated contexts (medical records, financial disclosures, court testimony, identity verification)—remain fully

in force. Narrative autonomy operates only in the space *outside* those duties: in casual self-presentation, in public-facing biographies, in social-media activity not directed at any specific counterparty in formation of a transactional or fiduciary relationship. A workable test: narrative autonomy attaches to self-presentations in non-transactional public contexts, where no specific counterparty is justified in detrimental reliance on the truth of the statement.

Public accountability erosion. Public figures—politicians, executives, holders of fiduciary or regulatory duties—are subject to public-interest scrutiny that depends on stable information about their conduct. The framework requires a public-figure exception, parallel to those embedded in defamation law (the *New York Times v. Sullivan* line) and right-to-be-forgotten doctrine (the *Google Spain* line). The category will remain difficult; that is a general feature of the legal landscape, not a specific objection.

Organised disinformation. Coordinated production of inconsistent narratives by state actors or organised interests is properly understood as disinformation, not privacy. The relevant tests concern target (oneself versus third parties), scale (individual versus coordinated), and purpose (resistance to profiling versus political manipulation). Narrative autonomy attaches to first-person self-presentation; disinformation operations involve third-party manipulation, scaled coordination, and political objectives.

Trust infrastructure erosion. At scale, narrative autonomy could erode the baseline trust on which public information environments depend. Two qualifications: it is unlikely to be practised at the scale the concern imagines, and even if widely adopted, the strategy targets a kind of trust already being undermined by generative AI itself. The relevant counterfactual is not "an open web of perfect reliability" but "an open web of degrading reliability."

Asymmetric impact on marginalised groups. Drawing on Eubanks (2018), narrative autonomy as a privacy strategy may be disproportionately accessible to those with the literacy, leisure, and platform familiarity to deploy it effectively. Marginalised groups—those most vulnerable to algorithmic harm—may be least able to engage in active narrative management, while remaining heavily exposed to algorithmic profile-based decisions. Narrative autonomy is not, on its own, a sufficient response to algorithmic injustice. Regulatory mechanisms should not be substituted by individual strategies; both are needed, and the regulatory ones are more important for marginalised populations.

7. Toward a New Governance Framework

Recognition of narrative autonomy in interpretive doctrine. Existing data protection laws need not be rewritten. What is required is interpretive recognition that publishing multiple self-narratives, absent fraud or defamation, is a protected exercise of expressive freedom. Platform terms of service requiring unitary "real identity" should be reviewed against this principle.

Profile stability as a regulatory object. High-stakes automated systems—credit scoring, hiring screens, judicial risk assessment, immigration decisions—should be required to disclose confidence estimates for their judgments about individuals. Where confidence is low, decisions based on those judgments should be subject to enhanced individual challenge rights. Profile stability is a measurable property; existing AI evaluation methods can be adapted to measure it.

Privacy-aware AI design. Designs that artificially restore stability through aggressive provenance tracking warrant scrutiny where individuals have a privacy interest in profile instability. Conversely, where downstream applications need reliability (medical, legal, journalistic), system design should support that reliability while honouring the privacy interest in ordinary-person profile uncertainty.

8. Conclusion

Privacy governance has, for two decades, been organised around the proposition that the threat is data existing and the response is data ceasing to exist. As generative AI takes over the role of public-facing information intermediary, this framing addresses a diminishing share of actual privacy harm. The new harm is profile integration, not data persistence; and the new resistance, accordingly, is profile destabilisation rather than data erasure.

This is a paradigm shift, in the modest sense that the questions one asks change. The older question—*can I have my data deleted?*—remains important. But it is increasingly accompanied by a different question: *can I refuse to be reduced to a stable profile?*

What is at stake, finally, is whether the institutional treatment of persons in the AI era can accommodate the basic fact that human lives are complex, contradictory, and not naturally compressible into single fluent representations. Erasure was the answer of the search-engine era. Narrative autonomy may be one of the answers of the LLM era.

References

- Bender, E. M., Gebru, T., McMillan-Major, A., & Shmitchell, S. (2021). On the dangers of stochastic parrots: Can language models be too big? In *Proceedings of the 2021 ACM Conference on Fairness, Accountability, and Transparency* (pp. 610–623). ACM.
- boyd, d., & Crawford, K. (2012). Critical questions for big data. *Information, Communication & Society*, 15(5), 662–679.
- Brunton, F., & Nissenbaum, H. (2015). *Obfuscation: A user's guide for privacy and protest*. MIT Press.
- Couldry, N., & Mejias, U. A. (2019). *The costs of connection*. Stanford University Press.
- Eubanks, V. (2018). *Automating inequality: How high-tech tools profile, police, and punish the poor*. St. Martin's Press.
- European Parliament and Council. (2016). Regulation (EU) 2016/679 (GDPR), Article 17. *Official Journal of the European Union*, L 119/1.
- Floridi, L. (2023). *The ethics of artificial intelligence*. Oxford University Press.
- Gerstgrasser, M., Schaeffer, R., Dey, A., et al. (2024). Is model collapse inevitable? Breaking the curse of recursion by accumulating real and synthetic data. *arXiv:2404.01413*.
- Goffman, E. (1959). *The presentation of self in everyday life*. Doubleday.
- Google Spain SL v. Agencia Española de Protección de Datos*, Case C-131/12, ECLI:EU:C:2014:317 (CJEU, May 13, 2014).
- Hartzog, W., & Selinger, E. (2013, January 17). Obscurity: A better way to think about your data than "privacy." *The Atlantic*.
- Hartzog, W., & Stutzman, F. (2013). The case for online obscurity. *California Law Review*, 101(1), 1–49.
- Jones, M. L. (2016). *Ctrl + Z: The right to be forgotten*. NYU Press.
- Mantelero, A. (2013). The EU proposal for a general data protection regulation and the roots of the "right to be forgotten." *Computer Law & Security Review*, 29(3), 229–235.
- Marwick, A. E., & boyd, d. (2011). I tweet honestly, I tweet passionately. *New Media & Society*, 13(1), 114–133.
- McClain, C., & Sidoti, O. (2025, June 25). *34% of U.S. adults have used ChatGPT*. Pew Research Center.
- Mittelstadt, B. D., Allo, P., Taddeo, M., Wachter, S., & Floridi, L. (2016). The ethics of algorithms: Mapping the debate. *Big Data & Society*, 3(2), 1–21.
- New York Times Co. v. Sullivan*, 376 U.S. 254 (1964).
- Nissenbaum, H. (2010). *Privacy in context*. Stanford Law Books.
- O'Neil, C. (2016). *Weapons of math destruction*. Crown.
- Pasquale, F. (2015). *The black box society*. Harvard University Press.
- Shumailov, I., Shumaylov, Z., Zhao, Y., Papernot, N., Anderson, R., & Gal, Y. (2024). AI models collapse when trained on recursively generated data. *Nature*, 631(8022), 755–759.
- Standing Committee of the National People's Congress of the People's Republic of China. (2021). *Personal Information Protection Law*, Article 47.
- van Dijck, J. (2014). Datafication, dataism and dataveillance. *Surveillance & Society*, 12(2), 197–208.
- Weidinger, L., Uesato, J., Rauh, M., et al. (2022). Taxonomy of risks posed by language models. In *Proceedings of the 2022 ACM Conference on Fairness, Accountability, and Transparency* (pp. 214–229). ACM.

Zuboff, S. (2019). *The age of surveillance capitalism*. PublicAffairs.